

# 05 - Fast Manipulations for “Bigger” Data

## R Workshop – Data wRestling

Adam Loy and Susan Vander Plas

Iowa State University

May 15, 2013

# Corn production

We have data on corn production in the Cornbelt from 1958–2011  
(Source: NASS Quick Stats)

```
corn <- read.csv("corn.csv")  
names(corn)
```

```
## [1] "Year"          "State"          "StateFIPS"      "County"  
## [5] "CountyCode"    "Crop"           "SubCrop"        "Measurement"  
## [9] "Value"
```

```
levels(corn$Measurement)
```

```
## [1] "ACRES HARVESTED" "ACRES PLANTED"  "PRODUCTION"  
## [4] "YIELD"
```

```
dim(corn)
```

```
## [1] 296377      9
```

## Corn production

Suppose we want to find the mean yield over the timespan for each county

```
plyr.agg <- ddply(corn, .(State, StateFIPS, County, CountyCode,  
  Measurement), summarise, means = mean(Value))
```

Ok, that was fast, but what if we were interested in summaries by year?

- ▶ Some counties reported irrigated and nonirrigated land separately
- ▶ We need to aggregate within Year
- ▶ We just went from aggregating within 1,010 groups to 51,830 groups
- ▶ Adding Year to the ddply statement took almost 13 minutes on my laptop

# Speeding things up

How can we speed up this aggregation?

- ▶ Manually subset
  1. Subset counties with numerous measurements
  2. Use `ddply` on the reduced set
- ▶ Find a faster [one-line](#) alternative

```
library(data.table)
```

```
help(data.table)
```

# data.table

## Pros

- ▶ Data tables are faster than data frames
  - ▶ What took 13 minutes with ddp1y took half a second with a data.table (not everything will be this big of an improvement)
- ▶ Data tables are also data frames (not everything breaks)
- ▶ Like the `**ply` statements, we can write compact and readable code
- ▶ Google uses data.tables (cool factor?)

## Cons

- ▶ Data tables are definitely harder to use... at first

# Getting started I

## Creating a data.table

### 1. Creating a data table

```
## start with a data frame
df <- data.frame(x = rep(c("a", "b", "c"), each = 3), y = c(1,
  3, 6), z = 1:9)

## convert to a data table
dt <- data.table(df)
```

### 2. Functions used with data frames still work, but might look a little different

```
head(dt)

##      x y z
## 1:  a 1 1
## 2:  a 3 2
## 3:  a 6 3
## 4:  b 1 4
## 5:  b 3 5
## 6:  b 6 6
```

# Getting started II

## Creating a data.table

### 3. Indexing is a bit different

```
## Wrong way
```

```
dt[, 1]
```

```
## [1] 1
```

```
dt[, "x"]
```

```
## [1] "x"
```

```
## Right way
```

```
dt[, x]
```

```
## [1] a a a b b b c c c
```

```
## Levels: a b c
```

```
dt[1:2, ]
```

```
##      x y z
```

```
## 1: a 1 1
```

```
## 2: a 3 2
```

# Getting started III

## Creating a data.table

But we can force data table to act more like the data frames

```
dt[, 1, with = FALSE]
```

```
##      x  
## 1: a  
## 2: a  
## 3: a  
## 4: b  
## 5: b  
## 6: b  
## 7: c  
## 8: c  
## 9: c
```



# Keys I

- ▶ Data frames have a single row name
- ▶ Data tables can have many names for a single row called a *key*
  - ▶ Each data table can only have 1 key
  - ▶ The data table is sorted by the key
- ▶ To see what keys are set we type

```
tables()
```

```
##      NAME NROW MB COLS  KEY  
## [1,] dt      9 1  x,y,z  
## Total: 1MB
```

# Keys II

Keys allow for easy subsetting

```
setkey(dt, x)
tables()
```

```
##      NAME NROW MB COLS  KEY
## [1,] dt      9 1  x,y,z x
## Total: 1MB
```

```
dt["b", ]
```

```
##      x y z
## 1: b 1 4
## 2: b 3 5
## 3: b 6 6
```

# Fast grouping I

Let's make a larger data table

```
grpsize <- ceiling(10e6 / 26^2)
DF <- data.frame(x = rep(LETTERS,each = 26 * grpsize),
                 y = rep(letters, each = grpsize),
                 v = runif(grpsize * 26^2),
                 stringsAsFactors = FALSE)
```

```
dim(DF)
```

```
## [1] 10000068      3
```

```
DT <- data.table(DF)
```

```
setkey(DT, x, y)
```

```
tables()
```

```
##      NAME      NROW  MB COLS  KEY
## [1,] dt           9   1 x,y,z x
## [2,] DT 10,000,068 229 x,y,v x,y
## Total: 230MB
```

# Fast grouping II

```
DT[, sum(v)]
```

```
## [1] 5e+06
```

- ▶ The second argument in `DT[i, j]` is used for aggregation
- ▶ You can put one or more expressions here
- ▶ To aggregate by group use `by`

```
head(DT[, sum(v), by = x])
```

```
##      x      V1  
## 1: A 192568  
## 2: B 192283  
## 3: C 192592  
## 4: D 192228  
## 5: E 192256  
## 6: F 192462
```

## Fast grouping III

To sum by both groups we simply add y to by

```
DT[, sum(v), by = "x,y"] ## no space!
```

```
##      x y    V1
##  1: A a 7382
##  2: A b 7384
##  3: A c 7366
##  4: A d 7413
##  5: A e 7453
##  ---
## 672: Z v 7425
## 673: Z w 7420
## 674: Z x 7401
## 675: Z y 7382
## 676: Z z 7381
```

## Your Turn

Return to the corn production data, and convert

```
# Don't run this!  
plyr.agg <- dply(corn, .(Year, State, StateFIPS, County, CountyCode,  
  Measurement), summarise, means = mean(Value))
```

into a data table expression.

## More complex queries

`DT[i, j, by]`

- `i` restrict attention to certain rows by stringing statements together using `&` (and); `|` (or)
- `j` use `list()` to string together multiple aggregation statements
- `by` use `list()` to include multiple groups or by putting everything in quotes with no spaces

## More complex queries I

We have data from the Cook's County Sheriff's Office that tracks location

```
load("crime.rda")
str(crime)
```

```
## 'data.frame': 19207 obs. of  11 variables:
## $ charges_citation      : chr  "720 ILCS 5 12-3.4(a)(2) [16145" "6
## $ race                  : chr  "WH" "LW" "BK" "BK" ...
## $ age_at_booking        : int   26 37 18 32 49 26 41 56 40 20 ...
## $ gender                : chr  "M" "M" "M" "F" ...
## $ booking_date          : Date, format: "2013-01-20" ...
## $ jail_id               : chr  "2013-0120171" "2013-0120170" "2013
## $ bail_status           : chr  NA NA NA NA ...
## $ housing_location      : chr  "05-" "05-" "05-L-2-2-1" "17-WR-N-A
## $ charges               : chr  NA NA NA NA ...
## $ bail_amount           : int   5000 10000 5000 50000 5000 5000 250
## $ discharge_date_earliest: Date, format: NA ...
```



## More complex queries II

```
crime <- as.data.table(crime)
```

```
## A quick table
```

```
crime[, .N, by = bail_status]
```

```
##      bail_status      N
## 1:             NA 10318
## 2:          NO BOND  8086
## 3: Bond in Process   801
## 4: 10000000.00      1
## 5: 25000000.00      1
```

```
crime[!is.na(bail_status), .N, by = bail_status]
```

```
##      bail_status      N
## 1:          NO BOND 8086
## 2: Bond in Process  801
## 3: 10000000.00      1
## 4: 25000000.00      1
```

## Your turn

Create a summary of the Cook County data by race, gender, and age at booking. Include summaries for

- ▶ count
- ▶ average, standard deviation, minimum, and maximum bail

Tip: exclude cases with NAs in the calculations